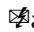


# **R Statistical Software in Action for Newcomers Data Analysis**

Arndt Leininger

 @a\_leininger

 arndt.leininger@fu-berlin.de

24 April 2018

## **Tuesday, 24 April 2018**

09:00h - 10:30h Data Analysis

*10:30h - 10:45h break*

10:45h - 12:15h Data visualization

12:15h - 12:30h Concluding remarks

## **Correlation and difference-in-means**

# Correlation

```
library(foreign)
d <- read.dta("../data/EUsuppDK.dta")

cor(d$age, d$left_right)
```

```
## [1] NA
```

```
cor(d$age, d$left_right, use = "complete.obs")
```

```
## [1] 0.07708635
```

# T-Test

```
t.test(left_right ~ sex, d)
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: left_right by sex
```

```
## t = 2.3134, df = 955.45, p-value = 0.02091
```

```
## alternative hypothesis: true difference in means is not equal
```

```
## 95 percent confidence interval:
```

```
## 0.04020583 0.48985106
```

```
## sample estimates:
```

```
## mean in group male mean in group female
```

```
## 5.568826 5.303797
```

- ▶ Missings (NA) are dropped automatically

# T-Test

```
t.test(left_right ~ sex, d, var.equal = T)
```

```
##
```

```
## Two Sample t-test
```

```
##
```

```
## data: left_right by sex
```

```
## t = 2.3064, df = 966, p-value = 0.0213
```

```
## alternative hypothesis: true difference in means is not equal
```

```
## 95 percent confidence interval:
```

```
## 0.03952888 0.49052801
```

```
## sample estimates:
```

```
## mean in group male mean in group female
```

```
## 5.568826 5.303797
```

## **Simple regression models**

# Regression

- ▶ Regression models are functions that are fed an equation and data
- ▶ Further options are possible but optional
- ▶ The dependent variable is separated by a tilde from the independent variables
- ▶ Equation:  $dv \sim iv$
- ▶ No need for \$ operator in the equation

```
lm(left_right ~ age, d)
```

```
##  
## Call:  
## lm(formula = left_right ~ age, data = d)  
##  
## Coefficients:  
## (Intercept)          age  
##      5.087151      0.007995
```



## **A complete overview of formulae in R**

https:

[//ww2.coastal.edu/kingw/statistics/R-tutorials/formulae.html](https://ww2.coastal.edu/kingw/statistics/R-tutorials/formulae.html)

# Model objects

- ▶ You can save the output of the `lm()` function just like with any other function.

```
m1 <- lm(left_right ~ age, d)
```

- ▶ `lm()` is for linear models, i.e. OLS

# Model objects

- ▶ Once you save an estimated model as object you can always access it to obtain model statistics.

```
summary(m1)  # estimation results
```

```
coef(m1)    # coefficients
```

```
vcov(m1)    # Variance-Covariance Matrix
```

```
predict(m1)  # Predicted values
```

```
resid(m1)   # Residuals
```

# Model objects

```
summary(m1)  # estimation results
```

```
##  
## Call:  
## lm(formula = left_right ~ age, data = d)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -4.687 -1.303 -0.351  1.515  4.649   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  5.087151   0.157297  32.341  <2e-16 ***  
## age          0.007995   0.003327   2.403   0.0164 *    
## ---  
## Signif. codes:  
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 1.787 on 966 degrees of freedom  
##    (33 observations deleted due to missingness)
```

# A little trick

```
summary(m <- lm(left_right ~ age, d))
```

```
##
## Call:
## lm(formula = left_right ~ age, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.687 -1.303 -0.351  1.515  4.649
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.087151   0.157297  32.341  <2e-16 ***
## age          0.007995   0.003327   2.403   0.0164 *
## ---
## Signif. codes:
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.787 on 966 degrees of freedom
## (33 observations deleted due to missingness)
```

# A little trick

```
# but  
summary(m = lm(left_right ~ age, d))
```

```
## Error in summary.lm(m = lm(left_right ~ age, d)): argument "o
```

# Model objects

```
coef(m1)  # coefficients
```

```
## (Intercept)          age  
## 5.087151366 0.007995059
```

# Model objects

```
vcov(m1)  # Variance-Covariance Matrix
```

```
##              (Intercept)              age
## (Intercept)  0.0247424372 -4.872123e-04
## age          -0.0004872123  1.106937e-05
```



# Model objects

```
predict(m1)  # Predicted values
```

##	1	2	3	4	5	6
##	5.414949	5.422944	5.215072	5.303018	5.279033	5.223067
##	7	8	9	10	11	12
##	5.231062	5.295023	5.295023	5.303018	5.231062	5.414949
##	13	14	15	16	17	18
##	5.311013	5.303018	5.279033	5.239057	5.319008	5.279033
##	19	20	21	22	23	24
##	5.430939	5.422944	5.327003	5.247053	5.438934	5.446929
##	25	26	27	28	29	30
##	5.430939	5.271038	5.255048	5.263043	5.454924	5.462919
##	31	32	33	34	35	36
##	5.470914	5.311013	5.319008	5.327003	5.438934	5.334998
##	37	38	39	40	41	42
##	5.334998	5.271038	5.223067	5.478909	5.446929	5.342993
##	43	44	45	46	47	48
##	5.454924	5.350988	5.358983	5.486904	5.462919	5.494899
##	49	50	51	52	53	54
##	5.279033	5.266079	5.249023	5.470914	5.502904	5.250988

# Model objects

```
resid(m1)  # Residuals
```

```
##           1           2           3           4           5
## -0.4149488 -0.4229438 -0.2150723  1.6969820  2.7209672
##           6           7           8           9          10
## -1.2230674 -0.2310624 -0.2950229 -0.2950229 -0.3030180
##          11          12          13          14          15
## -0.2310624  1.5850512  0.6889870  1.6969820 -3.2790328
##          16          17          18          19          20
##  1.7609425  2.6809919  2.7209672 -1.4309389 -0.4229438
##          21          22          23          24          25
## -1.3270031 -0.2470525  1.5610660 -0.4469290  1.5690611
##          26          27          28          29          30
## -3.2710377 -2.2550476 -3.2630427 -0.4549241  1.5370809
##          31          32          33          34          35
## -0.4709142 -0.3110130  1.6809919 -3.3270031  2.5610660
##          36          37          38          39          40
##  0.6650018 -0.3349982 -3.2710377 -0.2230674  2.5210907
##          41          42          43          44          45
##  1.5530710 -1.3429932  0.5450759 -0.3509883  2.6410166
```

# Model objects

```
class(m1)
```

```
## [1] "lm"
```

```
objects(m1)
```

```
## [1] "assign"      "call"        "coefficients"  
## [4] "df.residual" "effects"     "fitted.values"  
## [7] "model"      "na.action"   "qr"  
## [10] "rank"       "residuals"  "terms"  
## [13] "xlevels"
```

# Interactions

- ▶ Interactions can be specified as follows
  - ▶  $\text{var1} * \text{var2} = \text{var1} + \text{var2} + \text{var1}:\text{var2}$
  - ▶  $\text{var1}:\text{var2}$  is simply the interaction term

```
m_i <- lm(left_right ~ sex * age, d)
# or
m_i <- lm(left_right ~ sex + age + sex:age, d)
```

# Polynomials

```
lm(left_right ~ age + I(age^2), d)  # second-order polynomial  
lm(left_right ~ age + age^2, d)    # does not work
```

- ▶ ?I(): “Change the class of an object to indicate that it should be treated ‘as is.’”

## A little trick

- Put the assignment of a model to an object in a `summary()` call to assign and view results at the same time

```
summary(m2 <- lm(left_right ~ age + I(age^2), d))
```

```
##
## Call:
## lm(formula = left_right ~ age + I(age^2), data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.7388 -1.3154 -0.3393  1.5501  4.6607
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.3107878  0.3941342  13.475  <2e-16 ***
## age         -0.0029424  0.0179840  -0.164    0.870
## I(age^2)     0.0001153  0.0001863   0.619    0.536
## ---
## Signif. codes:
```

# Time-Series analysis

- ▶ Lags and leads can be easily included in time-series data
- ▶ For this make sure that the data are sorted on the time variable
  - ▶ Remember to use `order()` not `'sort()'` (which is for vectors) to sort a `data.frame`
  - ▶ Or use `dplyr`'s `arrange()`
- ▶ `lag()` and `lead()`
- ▶ For second lag: `lag(var, 2)`

```
summary(m_l <- lm(y ~ x + lag(x) + lag(x, 2), d))
```

# Logistic regression

- Generalized linear models, such as probit or logistic regression, are provided through the `glm()` function

```
d$wealthy <- ifelse(d$wealth == "++" | d$wealth == "+",  
  T, F)  
  
# Probit  
m_g <- glm(wealthy ~ age + sex, data = d)  
  
# Logistic  
m_g <- glm(wealthy ~ age + sex, family = binomial(), data = d)
```



# Predicted values

- ▶ `predict()` is a generic function to create predictions from various models

```
predicted_values <- predict(m1)
# but...
d$yhat <- predict(m1)
```

```
## Error in `$<-.data.frame`(`*tmp*`, yhat, value = structure(c(
```

```
# Error in `$<-.data.frame`(`*tmp*`, 'yhat', value =
# c(5.41494877919278, : replacement has 968 rows, data
# has 1001
```

# Predicted values

```
d$yhat <- predict(m1, newdata = d)  
# no error message because now predictions are also made  
# for deleted observations; these predictions are  
# obviously NA
```

# Predicted values

Or...

```
install.packages("broom")
```

```
library(broom)  
augment(m1) %>% head
```

```
##      .rownames left_right age  .fitted    .se.fit  
## 1           1           5  41 5.414949 0.05829785  
## 2           2           5  42 5.422944 0.05781844  
## 3           3           5  16 5.215072 0.10947787  
## 4           4           7  27 5.303018 0.08063835  
## 5           5           8  24 5.279033 0.08793296  
## 6           6           4  17 5.223067 0.10665959  
##      .resid      .hat    .sigma    .cooksd  
## 1 -0.4149488 0.001064565 1.787636 2.876892e-05  
## 2 -0.4229438 0.001047128 1.787634 2.939764e-05  
## 3 -0.2150723 0.003754221 1.787673 2.740268e-05  
## 4  1.6969820 0.002036810 1.786850 9.223870e-04  
## 5  2.7209672 0.002421980 1.785534 2.822023e-03
```

# Predicted values

```
df <- data.frame(age = min(d$age):max(d$age))
```

```
predict(m1, newdata = df)
```

##	1	2	3	4	5	6
##	5.207077	5.215072	5.223067	5.231062	5.239057	5.247053
##	7	8	9	10	11	12
##	5.255048	5.263043	5.271038	5.279033	5.287028	5.295023
##	13	14	15	16	17	18
##	5.303018	5.311013	5.319008	5.327003	5.334998	5.342993
##	19	20	21	22	23	24
##	5.350988	5.358983	5.366978	5.374973	5.382969	5.390964
##	25	26	27	28	29	30
##	5.398959	5.406954	5.414949	5.422944	5.430939	5.438934
##	31	32	33	34	35	36
##	5.446929	5.454924	5.462919	5.470914	5.478909	5.486904
##	37	38	39	40	41	42
##	5.494899	5.502894	5.510889	5.518885	5.526880	5.534875
##	43	44	45	46	47	48
##	5.542870	5.550865	5.558860	5.566855	5.574850	5.582845

## **Hands-on I**

# Hands-on I

[https://gitlab.com/arndtl/r\\_workshop](https://gitlab.com/arndtl/r_workshop)

## **Presenting Results with R**

# Lists

- ▶ Lists can contain any kind of objects of any type.
- ▶ Note: data.frames can also contain vectors of any of the three types but the vectors are forced to be of the same length.
- ▶ Example: One could have a list of differently sized vectors

```
v1 <- c(1, 2, 3)
v2 <- c("a", "b", "c", "d")
alist <- list(v1, v2)
alist
```

```
## [[1]]
## [1] 1 2 3
##
## [[2]]
## [1] "a" "b" "c" "d"
```



# Tables

Using packages such as `stargazer` or `texreg` we can create nice regression tables.

```
library(stargazer)  
library(texreg)
```

# Tables

Here's an example in tex using stargazer.

```
m1 <- lm(mpg ~ cyl, mtcars)
m2 <- lm(mpg ~ cyl + gear, mtcars)

stargazer(list(m1, m2), header = F, float = F,
            font.size = 'tiny', single.row = T)
```

Dependent variable:		
	mpg	
	(1)	(2)
cyl	-2.876*** (0.322)	-2.743*** (0.373)
gear		0.652 (0.904)
Constant	37.885*** (2.074)	34.659*** (4.937)
Observations	32	32
R <sup>2</sup>	0.726	0.731
Adjusted R <sup>2</sup>	0.717	0.712
Residual Std. Error	3.206 (df = 30)	3.232 (df = 29)
F Statistic	79.561*** (df = 1; 30)	39.404*** (df = 2; 29)

Note:

\* p<0.1; \*\* p<0.05; \*\*\* p<0.01

# Tables

stargazer provides tables in text, html and tex.

```
# output as text file  
stargazer(m1, type = "text", out = "tables/m1.txt")  
# output as html file which Word can read  
stargazer(m1, type = "html", out = "tables/m1.html")  
# output as tex, the default  
stargazer(m1, out = "tables/m1.tex")
```

# Screenreg

- texreg's screenreg() function is very useful to quickly view some models.

```
screenreg(list(m1, m2))
```

```
##
## =====
##               Model 1      Model 2
## -----
## (Intercept)  37.88 ***   34.66 ***
##              (2.07)      (4.94)
## cyl          -2.88 ***   -2.74 ***
##              (0.32)      (0.37)
## gear                             0.65
##                             (0.90)
## -----
## R^2           0.73        0.73
## Adj. R^2      0.72        0.71
## Num. obs.     32          32
## RMSE          3.21        3.23
```

# Tables

texreg provides tables html, tex and to screen.

```
# output as tex file  
texreg(m1, file = "tables/m1.tex")  
# output as html file which Word can read  
htmlreg(m1, file = "tables/m1.html")  
# output as tex, the default  
screenreg(m1)
```

# Regression tables for Word

## Via HTML

- ▶ Export to HTML using `htmlreg()` (package `texreg`) or `stargazer(..., type = 'html')`
- ▶ Then copy and paste to Word
- ▶ Or, better, link to the html file from within the Word Document
  - ▶ Word: Insert -> Object (dropdown) -> Text from File -> Insert (dropdown) -> Insert as link; hit F9 to refresh
  - ▶ LibreOffice Write: Insert -> Section -> Check option “Link” and choose document; to refresh: Edit -> Links -> click “Update”
  - ▶ <http://www.techrepublic.com/article/link-to-another-file-in-your-word-document/>

# Important packages for regression analysis

- ▶ `lmtest` for F-test and other tests
- ▶ `tseries` for time-series analysis
- ▶ `plm` for panel data analysis
- ▶ `lme4` for multilevel models
- ▶ `MatchIt` and `Matching` for matching

## **Hands-on II**



# Hands-on II

[https://gitlab.com/arndtl/r\\_workshop](https://gitlab.com/arndtl/r_workshop)

## **Further packages**

# ReporterRs

- ▶ ReporterRs is a package allows creation of entire(!) Word and Power Point documents
- ▶ this includes (regression) tables

```
install.packages("ReporterRs")
```

<http://davidgohel.github.io/ReporterRs/index.html>

# Regression tables with ReporterRs

```
library(ReporterRs)

# save the model in a data.frame
mdata <- as.data.frame(summary(m1)$coefficients)

# Define significance cutoffs
signif.codes <- cut(mdata[,4],
                    breaks = c( -Inf, 0.001, 0.01, 0.05, Inf),
                    labels= c("***", "**", "*", "" ) )

#format the values of coefficients, etc.
mdata[, 1:3] <- apply(mdata[, 1:3], 2, round, 2)
mdata[, 4] <- ifelse(mdata[, 4] < .001, "< 0.001",
                    round(mdata[, 4], 3))

# add signif codes to data
mdata$Signif = signif.codes
```

# Regression tables with ReporterRs

```
mdata
```

```
# create an empty FlexTable
```

```
coef_ft = FlexTable(data = mdata, add.rownames=TRUE,  
                    body.par.props = parRight(),  
                    header.text.props = textBold(),  
                    header.columns = T)
```

# Regression tables with ReporterRs

```
# format the table a bit  
coef_ft = setFlexTableBorders(coef_ft,  
                                inner.vertical = borderNone(),  
                                inner.horizontal = borderNone(),  
                                outer.vertical = borderNone(),  
                                outer.horizontal = borderSolid())
```

# Regression tables with ReporteRs

Now we'll save the table in a Word document.

```
# Create an empty document  
doc <- docx()  
  
# Add the regression table to the document  
doc <- addFlexTable(doc, coef_ft)  
  
# Save the document to the hard drive  
writeDoc(doc, file = "../tables/regtable.docx")
```